



COM106:

Introduction to

Databases

Further SQL – More
Operators and JOIN
Queries

Further SQL – More Operators and JOIN Queries

SQL provides a number of **Set Operators** to combine sets of rows returned by queries

We will look at **UNION**, **EXCEPT** and **INTERSECT**

UNION - Combines two query statements (tables) across **compatible** attributes

- Requires the *same number of attributes* in each query
- Each corresponding pair of attributes must be *defined on the same domain* (compatible)
- Duplicate values are eliminated

`SELECT <attribute_1>, < attribute_n>`

`FROM tables`

`[WHERE conditions]`

UNION


`SELECT <attribute_1>, < attribute_n>`

`FROM tables`

`[WHERE conditions];`

Gives Dataset 1

Gives Dataset 2



The **UNION** operator selects only **distinct values** by default.

To allow duplicate values, use the ALL keyword - **UNION ALL**

*The **UNION** operator returns records in the shaded area - records that exist in Dataset 1 or Dataset 2*

Further SQL – More Operators and JOIN Queries

INTERSECT - Returns values found in both queries

- Requires the *same number of attributes* in each query
- Each corresponding pair of attributes must be *defined on the same domain* (compatible)


```
SELECT <attribute_1>, .... < attribute_n>  
FROM tables  
[WHERE conditions]
```

INTERSECT

```
SELECT <attribute_1>, .... < attribute_n>  
FROM tables  
[WHERE conditions];
```

Gives Dataset 1

Gives Dataset 2



EXCEPT - Returns values found in the first query
not found in right query

- Requires the *same number of attributes* in each query
- Each corresponding pair of attributes must be *defined on the same domain* (compatible)

The **INTERSECT** operator returns records in the shaded area - records that exist in **both** Dataset 1 **and** Dataset 2

Further SQL – More Operators and JOIN Queries


```
SELECT <attribute_1>, .... < attribute_n>  
FROM tables  
[WHERE conditions]
```

EXCEPT

```
SELECT <attribute_1>, .... < attribute_n>  
FROM tables  
[WHERE conditions];
```

Gives Dataset 1

Gives Dataset 2



An Example of EXCEPT

A company holds details of its **SUPPLIERS** and its **CUSTOMERS**

SUPPLIER		
sup_no	s_name	country
s121	Big IT	Germany
s135	Top Design	Italy
s147	ALY Ltd	USA
s216	Microsoft	France
s227	Apple	Zambia
s251	PC Sport	Vietnam

CUSTOMER		
cus_no	c_name	contact
c164	Apple	Jobs
c045	IT r us	Sinclair
c237	Microsoft	Gates
c256	Myco	Jones
c290	MFN Ltd	Smith
c515	HP	Bloggs

Get the names of suppliers that are not also customers

s_name and c_name must have the same datatype

*The **EXCEPT** operator returns records in the shaded area - records that exist in Dataset 1 and not in Dataset 2*

```
SELECT s_name  
FROM SUPPLIER  
EXCEPT  
SELECT c_name  
FROM CUSTOMER
```

s_name
Big IT
Top Design
ALY Ltd
PC Sport

Microsoft and Apple from SUPPLIER are not in the results because they occur in CUSTOMER

Further SQL – More Operators and JOIN Queries

An Example of INTERSECT

SUPPLIER

sup_no	s_name	country
s121	Big IT	Germany
s135	Top Design	Italy
s147	ALY Ltd	USA
s216	Microsoft	France
s227	Apple	Zambia
s251	PC Sport	Vietnam

CUSTOMER

cus_no	c_name	contact
c164	Apple	Jobs
c045	IT r us	Sinclair
c237	Microsoft	Gates
c256	Myco	Jones
c290	MFN Ltd	Smith
c515	HP	Bloggs

**Get the names of suppliers that
are also customers**

*s_name and c_name
must have the same
datatype*

SELECT s_name
FROM SUPPLIER

INTERSECT

SELECT c_name
FROM CUSTOMER

s_name
Microsoft
Apple

*Microsoft and Apple are in
both SUPPLIER and CUSTOMER*

An Example of UNION

SUPPLIER

sup_no	s_name	country
s121	Big IT	Germany
s135	Top Design	Italy
s147	ALY Ltd	USA
s216	Microsoft	France
s227	Apple	Zambia
s251	PC Sport	Vietnam

CUSTOMER

cus_no	c_name	contact
c164	Apple	Jobs
c045	IT r us	Sinclair
c237	Microsoft	Gates
c256	Myco	Jones
c290	MFN Ltd	Smith
c515	HP	Bloggs

**Get the names of all
suppliers and customers**

*s_name and c_name
must have the same
datatype*

SELECT s_name
FROM SUPPLIER

UNION

SELECT c_name
FROM CUSTOMER

s_name
Big IT
Top Design
ALY Ltd
Microsoft
Apple
PC Sport
IT r us
Myco
MFN Ltd
HP

*All records in
SUPPLIER and
CUSTOMER*

Further SQL – More Operators and JOIN Queries

SQL also provides a number of **Logical Operators** used to test for the truth of some condition. We have already seen **AND, OR, BETWEEN, NOT** and **IN** – (*there are a number of others*)

The LIKE Operator

LIKE allows the use of **wildcards** to perform **pattern matching** on an attribute in a query. It can be used in the **WHERE** clause of a **SELECT, INSERT, UPDATE, or DELETE** statement.

- % wildcard representing any string of characters
- _ wildcard representing a single character
- [] looks for any match within a range of characters e.g. [a-m]
- [^] looks for any match not in the specified range e.g. [^a-m]

Some Examples:

Get the names of employees starting with 'S'

```
SELECT ename  
FROM EMPLOYEE  
WHERE ename LIKE 'S%';
```



ename
Smith
Smart

Get the names of employees containing the string 'row' anywhere

```
SELECT ename  
FROM EMPLOYEE  
WHERE ename LIKE '%row%';
```



ename
Brown
Thrower

EMPLOYEE

enum	ename	salary	floor
852341	Smith	15000	1
852358	Smart	19000	3
852407	Brown	16000	3
852455	Bruce	25100	2
852491	Thrower	30500	1
852514	Dale	11650	2
852530	Dole	26980	4

Further SQL – More Operators and JOIN Queries

Get the names of employees matching ‘D?le’

```
SELECT ename  
FROM EMPLOYEE  
WHERE ename LIKE 'D_le';
```



ename
Dale
Dole

Get the names of employees with any characters from n to r

```
SELECT ename  
FROM EMPLOYEE  
WHERE ename LIKE '%[n-r]%' ;
```



ename
Smart
Brown
Bruce
Thrower

The IS NULL Operator

NULL is a special marker used in SQL to indicate that a data value **does not exist**

NULL is not a value – it is a **state** indicating ‘unknown’ or the **absence of a value**.

SQL provides two special Null-specific comparison **predicates**:

IS NULL – tests whether an attribute has a value and returns **TRUE** if it is NULL

IS NOT NULL – tests whether an attribute has a value and returns **TRUE** if it does

e.g., `SELECT attribute_a`

`FROM A_TABLE`

`WHERE attribute_b IS NULL`

EMPLOYEE

enum	ename	salary	floor
852341	Smith	15000	1
852358	Smart	19000	3
852407	Brown	16000	3
852455	Bruce	25100	2
852491	Thrower	30500	1
852514	Dale	11650	2
852530	Dole	26980	4

Not the same as a space (' ') or a zero



Further SQL – More Operators and JOIN Queries

JOIN Fundamentals

A **JOIN** is a means for combining columns from two or more tables by using values common to each.

There are a number of different ways in which two tables can be joined - ANSI-standard SQL specifies five types of join: **INNER**, **LEFT OUTER**, **RIGHT OUTER**, **FULL OUTER** and **CROSS**.

As a special case, a table can join to itself in a **SELF-JOIN**.

The JOIN condition is specified in the FROM clause as:

*We've already seen examples of **INNER** JOINS and **SELF** JOINS*

*We will not consider **CROSS** JOINS here*

*FROM table1 **join_type** table2 [**ON** (*join_condition*)]*

The **join_condition** defines the way two tables are related in a query by specifying:

- The column from each table to be used for the join (typically, *but not necessarily*, on a PK/FK match).
- A logical operator (for example, = or <>, etc) to be used in comparing values in the columns.

*'=' is most commonly used – also known as an **equi-join***

INNER JOINS **only** can also be specified in the WHERE clause

Mark 1 - *implicit join notation*

```
SELECT ename, role  
FROM EMPLOYEE, WORKS_ON  
WHERE EMPLOYEE.enum = WORKS_ON.enum  
AND floor = 1;
```

Mark 2 - *explicit join notation*

```
SELECT ename, role  
FROM EMPLOYEE INNER JOIN WORKS_ON  
ON EMPLOYEE.enum = WORKS_ON.enum  
WHERE floor = 1;
```

Further SQL – More Operators and JOIN Queries


Explicit join notation is considered **best practice** since the join conditions are separated from any other conditions in the WHERE clause.

All JOIN queries create a **new result table** by combining column values of two tables (TABLE_A and TABLE_B) based upon the **join condition**. ← *Also known as the join predicate*

Each row of TABLE_A is compared with each row of TABLE_B to find all pairs of rows which satisfy the **join predicate**.


The **join type** specifies how the tables are to be combined (*as in the Venn Diagrams below*):

TABLE_A **INNER JOIN** TABLE_B



INNER JOIN - Select all records from TABLE_A and TABLE_B, where the join condition is met.


TABLE_A **LEFT OUTER JOIN** TABLE_B



LEFT OUTER JOIN - Select all records from TABLE_A, along with records from TABLE_B for which the join condition is met (if at all).

Further SQL – More Operators and JOIN Queries

TABLE_A **RIGHT OUTER JOIN** TABLE_B



RIGHT OUTER JOIN - Select all records from TABLE_B, along with records from TABLE_A for which the join condition is met (if at all).

Consider the following schema:

CLUB (club_id, club_name)


MEMBER (member_id, member_name, club_id*)

With sample data as shown:

We will look at the result of the different **join types** based on the **join condition**:-

MEMBER.club_id = CLUB.club_id

TABLE_A **FULL OUTER JOIN** TABLE_B



FULL OUTER JOIN - Select all records from TABLE_A and TABLE_B, regardless of whether the join condition is met or not.

MEMBER

member_id	member_name	club_id
M20	Smith	C3
M21	Jones	Null
M22	White	C1
M23	Black	C3
M24	Green	C1
M25	Brown	Null

CLUB

club_id	club_name
C1	Tennis
C2	Archery
C3	Judo

Further SQL – More Operators and JOIN Queries

INNER JOIN

```
SELECT *  
FROM MEMBER INNER JOIN CLUB  
ON MEMBER.club_id = CLUB.club_id;
```

Note: Members "Jones" and "Brown" and the "Archery" club do not appear in the results. Neither of these has any matching rows in the other respective table.

LEFT OUTER JOIN

```
SELECT *  
FROM MEMBER LEFT OUTER JOIN CLUB  
ON MEMBER.club_id = CLUB.club_id;
```

Note: A **LEFT JOIN** returns all the values from the left table, plus matched values from the right table or **NULL** in case of no matching join predicate.

RESULT TABLE

member_id	member_name	club_id	club_name
M20	Smith	C3	Judo
M22	White	C1	Tennis
M23	Black	C3	Judo
M24	Green	C1	Tennis

RESULT TABLE

member_id	member_name	club_id	club_name
M20	Smith	C3	Judo
M21	Jones	Null	Null
M22	White	C1	Tennis
M23	Black	C3	Judo
M24	Green	C1	Tennis
M25	Brown	Null	Null

Further SQL – More Operators and JOIN Queries

RIGHT OUTER JOIN

```
SELECT *  
FROM MEMBER RIGHT OUTER JOIN CLUB  
ON MEMBER.club_id = CLUB.club_id;
```

Note: A **RIGHT JOIN** is similar to a LEFT JOIN, but with the treatment of the tables reversed. Returns all the values from the right table, plus matched values from the left table or *NULL* in case of no matching join predicate.

Notice- MEMBER **RIGHT OUTER JOIN** CLUB gives the same result as CLUB **LEFT OUTER JOIN** MEMBER

RESULT TABLE

member_id	member_name	club_id	club_name
M20	Smith	C3	Judo
M22	White	C1	Tennis
M23	Black	C3	Judo
M24	Green	C1	Tennis
Null	Null	C2	Archery

The order of the tables is important

FULL OUTER JOIN

```
SELECT *  
FROM MEMBER FULL OUTER JOIN CLUB  
ON MEMBER.club_id = CLUB.club_id;
```

Note: A **FULL JOIN** combines the results of both left and right outer joins.

The joined table will contain all records from both tables, with *NULLs* for missing matches on either side.

Some RDBMS do not support FULL JOINS, eg., MySQL

RESULT TABLE

member_id	member_name	club_id	club_name
M20	Smith	C3	Judo
M21	Jones	Null	Null
M22	White	C1	Tennis
M23	Black	C3	Judo
M24	Green	C1	Tennis
M25	Brown	Null	Null
Null	Null	C2	Archery

Further SQL – More Operators and JOIN Queries

Some Examples of Use

List all members and the clubs to which they belong.

Using an INNER JOIN

```
SELECT *
FROM MEMBER INNER JOIN CLUB
ON MEMBER.club_id = CLUB.club_id;
```

member_id	member_name	club_id	club_name
M20	Smith	C3	Judo
M22	White	C1	Tennis
M23	Black	C3	Judo
M24	Green	C1	Tennis

Using a LEFT JOIN

```
SELECT *
FROM MEMBER LEFT OUTER JOIN CLUB
ON MEMBER.club_id = CLUB.club_id;
```

member_id	member_name	club_id	club_name
M20	Smith	C3	Judo
M21	Jones	Null	Null
M22	White	C1	Tennis
M23	Black	C3	Judo
M24	Green	C1	Tennis
M25	Brown	Null	Null

Notice:- The **INNER JOIN** excludes members that don't have a club

The **LEFT JOIN** lists all members, *including* those that don't have a club

INNER JOIN is the most commonly used join, but should be **treated with care**, especially if the attribute on which the tables are joined can contain **NULLs**.

Further SQL – More Operators and JOIN Queries

Get the members that do not belong to any club

Using a LEFT JOIN

```
SELECT member_id, member_name  
FROM MEMBER LEFT OUTER JOIN CLUB  
ON MEMBER.club_id = CLUB.club_id  
WHERE CLUB.club_id IS NULL;
```

member_id	member_name
M21	Jones
M25	Brown

Get the clubs that do not have any members

Using a RIGHT JOIN

```
SELECT C.club_id, club_name  
FROM MEMBER AS M RIGHT OUTER JOIN CLUB AS C  
ON M.club_id = C.club_id  
WHERE M.member_id IS NULL;
```

club_id	club_name
C2	Archery

Notice the use of aliases:
MEMBER AS M,
CLUB AS C

Compare INTERSECT and INNER JOIN

Some queries can be answered using either an **INTERSECT** or an **INNER JOIN** approach

e.g., Rewrite the **INTERSECT** query on slide 5 using an **INNER JOIN**

Can the **INNER JOIN** query on slide 13 be rewritten using **INTERSECT?** **NO**